

---

# Towards More Explainable Multimodal Machine Learning

---

Yuanxin (Michael) Wang<sup>1</sup> Jiaxin (Kelly) Shi<sup>1</sup>

## Abstract

Deep neural networks have been proven to be both effective and efficient in many tasks, however, the complex model architecture makes them difficult for us to understand their decision-making process. As human behaviors are intrinsically multimodal, multimodal deep neural networks have achieved better performance than unimodal baselines in most cases. With growing applications of multimodal deep neural networks based Artificial Intelligence, the interplay between heterogeneous and high-dimensional modalities makes explainability a key concern of multimodal neural networks, as its application in areas like healthcare requires trustworthy decisions. This paper introduces a gradient-based explanatory method that visualizes the contribution of features to both the cross-modal interaction and classification result of any black-box multimodal neural networks. This paper serves as a midterm progress report for our advanced multimodal machine learning class.

## 1. Introduction

Deep neural networks have made numerous breakthroughs in creating applications for different modalities including vision (Dosovitskiy et al., 2021) (He et al., 2021), language (Devlin et al., 2019) (Brown et al., 2020), and speech (Miao et al., 2015) (Rao et al., 2017). However, these complex models are blamed for a lack of trust and transparency given their intrinsic black-box nature. Correspondingly, the idea of explainable AI (XAI) has been adopted to probe the inner mechanisms of these unimodal architectures (Tjoa & Guan, 2021). Among a wide range of XAI approaches, the gradient-based explanatory method has achieved significant progress in understanding how each segment in the input contributes to the prediction results.

The design of explainable machine learning models is expected to be even more challenging when we move from unimodal to multimodal scenarios. Despite the recent increasing popularity in different domains (Baltruaitis et al., 2019), one key challenge in multimodal machine learning is to explain cross-modal interactions. Cross-modal feature extraction allows for richer representations, thus improving

accuracy. However, the interplay between heterogeneous and high-dimensional modalities makes explainability a key concern of multimodal neural networks, as its application in areas like healthcare requires trustworthy decisions. Moreover, cross-modal interactions occur at multi-stages during the learning process, which makes it harder for us to measure how much cross-modal interactions are modeled. Specifically, the following questions should be answered to better explain cross-modal interaction: (1) what is the contribution of each modality to the prediction? (2) which input modality is dominant? (3) how to identify important input segments that facilitate the interaction between modalities?

Noticing that the aforementioned gradient-based explanatory method was mostly applied to the unimodal cases and has not been explored in the space of cross-modal interaction, we believe that the marriage between these two fields becomes a natural next step for both multimodal machine learning and XAI communities. In this work, we propose a novel gradient-based paradigm to uncover the extent of cross-modal interaction in a multimodal classification model and what role each input segment (e.g., image patches, language tokens, and audio clips) plays to enable such interaction. During the back-propagation on a trained multimodal model, the gradients flow back from one specific prediction label, through all intermediate representations and layers, and finally towards all input segments from different modalities. By completely or partially perturbing the input modalities, various cross-modal interaction patterns can be observed through the magnitude changes in the gradients from different layers. To the best of our knowledge, this paper is the first to explore the correlation between gradient variations and cross-modal interaction. Our code will be available at [github.com/MichaelYxWang/CrossExplain](https://github.com/MichaelYxWang/CrossExplain).

## 2. Related Work

### 2.1. Explainable AI (XAI)

Artificial Intelligence systems based on deep neural networks are powerful in learning in many domains. Although deep neural networks have been proven to be effective and efficient for both unimodal and multimodal tasks, the complex model architecture and hidden layer processing make them difficult for us to understand their internal states and decision-making process. There are three major types of

XAI methods: intrinsic, distillation and visualization.

Intrinsic methods can be achieved through either attention mechanisms or joint training. Attention mechanisms (Vaswani et al., 2017a) focus on regions that are important for making predictions and can make the model inherently self-explainable. Jointly training the model for prediction and explanation can also make it intrinsically self-explanatory. For example, (Kanehira et al., 2019) jointly trains a classification module and an explanation module for predicting counterfactuality.

Distillation methods can be done through local approximation and model translation. LIME (Ribeiro et al., 2016) uses a simple and interpretable model to approximate the black-box model locally. It is model-agnostic, and it performs perturbations around a particular prediction and sees how the predictions change. The model translation method builds a surrogate model to simulate the original model. The surrogate model is usually inherently explainable. For example, (Kaya et al., 2017) uses decision trees to approximate a black-box model.

Visualization methods can be categorized into two types: backpropagation and perturbation methods. Backpropagation methods compute the feature relevance based on the gradients passed through the network. Some existing backpropagation works are Grad-CAM (Selvaraju et al., 2016), DeepLIFT (Shrikumar et al., 2019), and layer-wise relevance propagation (Binder et al., 2016). The perturbation method alters the input and compares the output with original and modified input to identify sensitive features for prediction (Zeiler & Fergus, 2013).

## 2.2. Visual Explanation

Understanding what regions contribute most to the final class label prediction is key to explainable computer vision models. Apart from directly observing intermediate convolutional activation maps, gradient-based methods such as Grad-CAM (Selvaraju et al., 2016) and class saliency map visualization (Simonyan et al., 2014), where the gradient of each raw pixel for a specific class label is computed and the resulting heatmap is used to show the relevance of each region. Most gradient-based methods can be universally applied to any Convolutional neural network architecture. Extensions to Grad-CAM including XGradCAM (Fu et al., 2020), EigenGradCAM (Muhammad & Yeasin, 2020), and LayerCAM (Jiang et al., 2021) have continuously increased the dimensions of computer vision explainability as well.

## 2.3. Textual Explanation

Compared to the direct computation of the raw pixel-level gradient to understand computer vision outputs, explaining word/token importance in natural language processing is a

different story. The mismatch between words represented as continuous embedding vectors and word-level importance scores represented as scalar values has motivated the sensitivity analysis (SA-based) method. SA-based methods operate directly on either the raw value (Nguyen, 2018), L1-norm (Li et al., 2016), or L2-norm (Arras et al., 2016) of the gradient vectors for each word and use the variants of these gradient vectors as importance scores. The limitation of the SA-based method is that it can only measure the absolute value of word importance instead of distinguishing between positive and negative effects. Gradient  $\times$  Input (GI) method (Denil et al., 2014) and its variants (Arras et al., 2019) (Ancona et al., 2018), where the dot product between the word vector and the gradient vector is performed to compute the word importance score, are proposed to solve this issue. In addition, layer-wise relevance propagation, where the local relevance redistribution of each input is set proportionally to its contribution in the forward pass, is applied in natural language processing to tell the differences between words that are used to support or oppose the classification decisions (Arras et al., 2016).

However, researchers have also noticed that gradient-based explanatory methods in natural language processing are not always reliable (Wang et al., 2020): gradients can be manipulated without affecting the prediction of the model in adversarial settings and can lead to completely unreasonable word importance scores (e.g., focusing only on stop words). Va NLP explainability can also be tackled by non-gradient-based methods effectively. Inspired by how convolutional neural networks can be used to represent sentences (Kim, 2014), word importance scores can be also obtained by performing average pooling on each row of the sentence matrix and summing up the pooled vectors across different feature maps for each word (Lee et al., 2018).

With the prevalence of BERT-family models, (Vaswani et al., 2017b) (Devlin et al., 2019), explainable transformer architectures have aroused increasing interest from the research community. In addition to attention head inspection and representation latent space visualization methods (Coenen et al., 2019) (Rogers et al., 2020), direct analysis of per-token local relevance can be achieved by propagating relevance scores based on the deep Taylor Decomposition principle through attention layers and skip connections (Chefer et al., 2021).

## 2.4. Multimodal Explanation

A growing number of deep neural network-based models have made significant progress on multimodal tasks based on several modalities such as vision, textual data, and language. Several studies have applied explainability methods to determine the relative importance of each modality in making decisions on multimodal datasets. We could categorize them by the stage explanation modeling is introduced

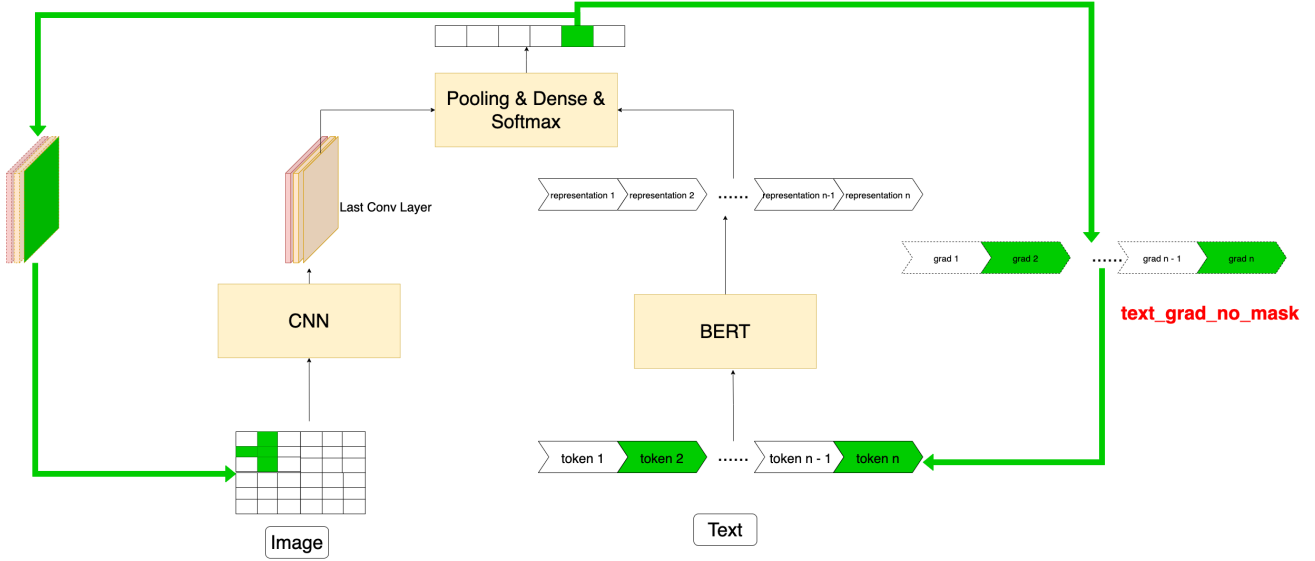


Figure 1. Forward and backward pass two full modalities through the model and get `text_grad_no_mask`.

to the deep neural networks: during modeling and post-modeling phase.

During the modeling phase, models are inherently developed to provide explanations by employing the intrinsic method. For generating a multimodal explanation for VQA, (Wu & Mooney, 2018) uses the model-agnostic explainer LIME (Ribeiro et al., 2016) to determine the segmented objects that influenced the decision the most; then it learns to embed the question, answer, and the VQA-attended features to generate textual explanations and measure how well the object referenced in the generated explanation matches the segments highlighted by LIME. There are also hybrid models that joint prediction and explanation. Contextual Explanation Networks (Al-Shedivat et al., 2017) is a class of probabilistic models that learns to predict by generating and leveraging intermediate context-specific explanations. Self-explanatory Neural Networks (Alvarez-Melis & Jaakkola, 2018) consists of three components: a concept encoder that transforms the input into a small set of interpretable basis features, an input-dependent parameterizer that generates relevance scores, and an aggregation function that makes predictions. Its robustness loss on the parameterizer encourages the full model to behave as a linear function locally, which yields an immediate interpretation of both concepts and relevance. For giving explanations in the post-hoc phase, the explanation method is implemented after the model is trained through backpropagation methods (Selvaraju et al., 2016) and proxy models (Kaya et al., 2017). These methods are model-agnostic thus can be applied to any trained models to improve their explainability.

### 3. Proposed Method

Given a trained multimodal classification model, a post-hoc gradient-based explanation method is developed to investigate which segments of the input contribute significantly to the final class prediction and which segments contributing more to the cross-modal interaction. To identify such segments for each input example, four rounds of forward-backward passes with different input perturbations are needed. For simplicity, our illustrated example uses a multimodal classification model that takes one image and one piece of text as input, and outputs a Softmax probability across various classes.

In the first round, as shown in Figure 1, both text and image modalities are forward passed through the model and the vector containing the gradients for each text token representation is denoted as `text_grad_no_mask`. This gradient vector is expected to identify the tokens (e.g., token 2 and token n in Figure 1) that play a more important role in predicting a specific class label compared to other tokens.

In the second round illustrated in Figure 5, we mask the image completely and only the text modality is forward passed. We denote the same gradient vector for all token representations as `text_grad_mask_all`. We compare the norms between this new vector with `text_grad_no_mask` to identify whether the model is learning from a dominant modality or from cross-modal interaction. If there is no change in the predicted class and no significant difference between the norms, it implies that the image modality is not contributing and the text modality is dominant; otherwise, there are two possible explanations: (1) the image modal-

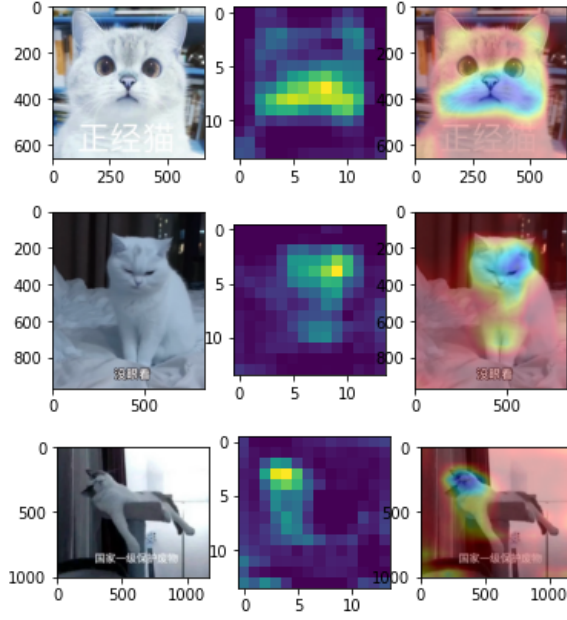


Figure 2. Unimodal results for image classification Grad-CAM visualization

ity is dominant (2) cross-modal interaction is necessary for this prediction. Explanation (1) can be relatively straightforward to examine by simply masking the text completely and observing the gradient norm changes in the image side. Assuming that we confirm that cross-modal interaction actually happened during the prediction, then intuitively the tokens that receive the largest gradients in this vector (e.g., token  $n-1$  and token  $n$  in Figure 5) should be regarded as bimodally-important.

To make this assumption more convincing, the third round of forward-backward pass is performed with the text modality completely masked to get the image gradient matrix denoted as *img\_grad\_mask\_all* and a fourth round is carried out with only the pre-assumed bimodally-important tokens (e.g., token  $n-1$  and token  $n$  in Figure 5) masked to get the image gradient matrix denoted as *img\_grad\_mask\_partial*, which are illustrated in Figure 7 and Figure ??, respectively. If these two matrices are similar to each other, we can conclude that the pre-assumed bimodally-important tokens play such a key role in the bimodal interaction that masking them only takes the same effect as masking the whole text modality.

By repeating a similar process on the image end, we can also find the image regions that are important for the prediction of a specific label as well as cross-modal interaction. Apart from masking input modalities, other modality-specific perturbation methods such as re-ordering the tokens in the text and adding Gaussian noise to the image pixels, can also be integrated into the pipeline.

## 4. Experimental Setup

### 4.1. Dataset

The [Hateful Memes dataset](#) contains 10,000+ new multimodal examples created by Facebook AI. The Hateful Memes dataset consists of more than 10,000 newly created examples of multimodal content. The memes were selected in such a way that strictly unimodal classifiers would struggle to classify them correctly (Figure 3). The dataset also contains multimodal memes that are similar to hateful examples but are actually harmless in order to help researchers address potential biases in classification systems and build systems that avoid false positives.

### 4.2. Baseline Models

For our project, we will analyze gradients of several baselines of the Hateful Memes dataset. Compared to human performance on this dataset being 84.70%, unimodal baselines such as Image-grid (Gao et al., 2020) and Text BERT (Devlin et al., 2019) can only achieve 52.0% and 59.2% respectively, while multimodal baselines such as ViLBERT (Lu et al., 2019) can achieve 61.1% test accuracy on the Hateful Memes dataset.

## 5. Current Results and Discussion

Building explainable multimodal models are generally more difficult than unimodal ones due to the added complexity from more modalities. Therefore, it is crucial to test the gradient-based method for unimodal cases first. Currently, we have implemented gradient-based explanatory methods for two unimodal models: a VGG-based image classification model and a BERT-based sentiment classification model.

For the vision side, we implement the Grad-CAM algorithm as illustrated in Figure 4. The gradients are pooled across each channel of the last convolutional feature map to several scalar values, and a weighted sum is performed between these pooled scalar values and the activation maps from the corresponding channel in the last convolutional layer to generate a class saliency heatmap. Initial results, as shown in Figure 2, show that the model can successfully classify cat images and highlight important regions (e.g., cat faces and mouths) to justify their predictions.

For the language side, we incorporate a gradient-based idea that is similar to Grad-CAM, where the gradients are back-propagated to the contextualized representations of each token and then element-wise multiplied with the representation vectors themselves. We use the norm of the resulting vector product to illustrate token importance to the final class prediction. However, we find that this method does not work well since the contributions are usually attributed randomly to  $\langle PAD \rangle$  and  $\langle SEP \rangle$  tokens. More work is needed





Figure 3. In each of these sample memes, the text phrase and the image are innocuous when considered by themselves. The semantic content of the meme becomes mean only when the text phrase and image are considered together.

to investigate the reason behind this.

The implementation of gradient-based explanatory method for the multimodal case is simply the aggregation of the unimodal cases. Currently, we are blocked by the inferior performance of the language side; but once this is resolved, we can start our cross-modal interaction probing iterations, as discussed in the previous sections.

## 6. Future Work

For the next steps, we plan to investigate the root cause of the poor performance for gradient-based method for languages and finish implementing the gradient-based method for a baseline vision-language multimodal classification model. Once the experimental pipeline is ready, we plan to run the perturbation cycles on different inputs and observe different patterns of cross-modal interaction. For the same input image-text pair, it would be interesting to see the differences between the baseline model and the current state-of-the-art models in terms of the input segments they regard as bimodally-important.

In the long term, we plan to design more robust quantitative evaluation metrics for our paradigm. In addition, we plan to extend our work from image-text models to video-text models and answer more complex multimodal explanatory questions such as (1) Which chunk of the video contributes most to the bimodal or trimodal interactions? (2) If conversations are involved in the video, which speaker turns contributes most to the interactions? (3) If multi-party conversations are involved in the video, which party contributes most to the interactions? We believe that our gradient-based framework is flexible for the extension to more modalities with temporal complexity.

## References

- Al-Shedivat, M., Dubey, A., and Xing, E. P. Contextual explanation networks. *CoRR*, abs/1705.10301, 2017. URL <http://arxiv.org/abs/1705.10301>.
- Alvarez-Melis, D. and Jaakkola, T. S. Towards robust interpretability with self-explaining neural networks. *CoRR*, abs/1806.07538, 2018. URL <http://arxiv.org/abs/1806.07538>.
- Ancona, M., Ceolini, E., Öztireli, C., and Gross, M. Towards better understanding of gradient-based attribution methods for deep neural networks. 01 2018.
- Arras, L., Horn, F., Montavon, G., Müller, K.-R., and Samek, W. Explaining predictions of non-linear classifiers in nlp. *ArXiv*, abs/1606.07298, 2016.
- Arras, L., Osman, A., Müller, K.-R., and Samek, W. Evaluating recurrent neural network explanations. In *Black-boxNLP@ACL*, 2019.
- Baltruaitis, T., Ahuja, C., and Morency, L.-P. Multimodal machine learning: A survey and taxonomy. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 41: 423–443, 2019.
- Binder, A., Montavon, G., Bach, S., Müller, K.-R., and Samek, W. Layer-wise relevance propagation for neural networks with local renormalization layers, 2016.
- Brown, T. B., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., Neelakantan, A., Shyam, P., Sastry, G., Askell, A., Agarwal, S., Herbert-Voss, A., Krueger, G., Henighan, T. J., Child, R., Ramesh, A., Ziegler, D. M., Wu, J., Winter, C., Hesse, C., Chen, M., Sigler, E., Litwin, M., Gray, S., Chess, B., Clark, J., Berner, C., McCandlish,

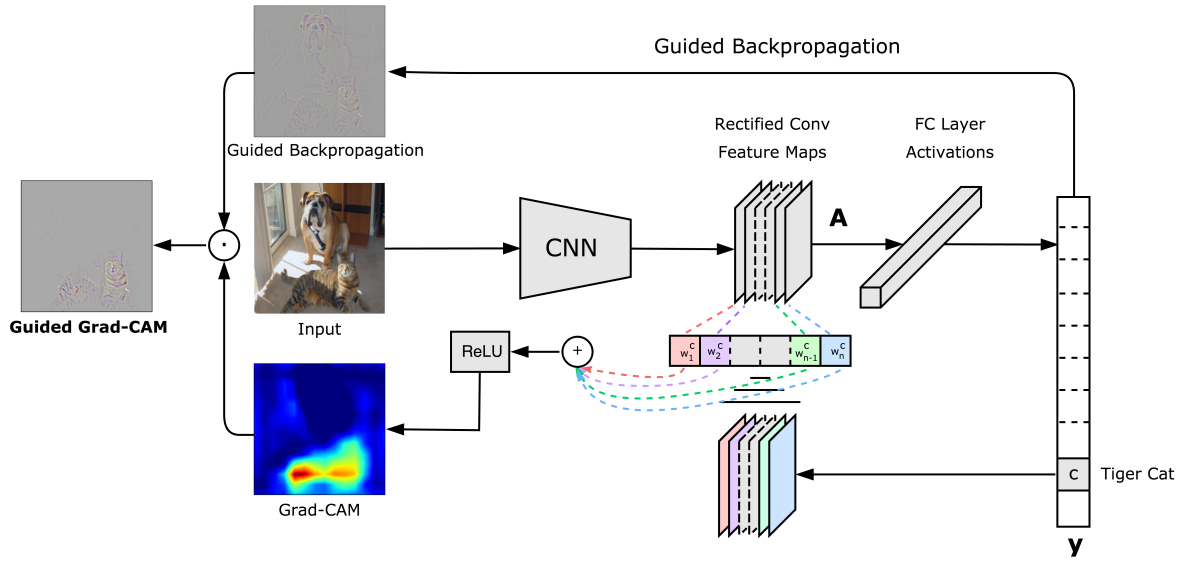


Figure 4. Grad-CAM illustration (Selvaraju et al., 2016)

- S., Radford, A., Sutskever, I., and Amodei, D. Language models are few-shot learners. *ArXiv*, abs/2005.14165, 2020.
- Chefer, H., Gur, S., and Wolf, L. Transformer interpretability beyond attention visualization. *2021 IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 782–791, 2021.
- Coenen, A., Reif, E., Yuan, A., Kim, B., Pearce, A., Viégas, F. B., and Wattenberg, M. Visualizing and measuring the geometry of bert. In *NeurIPS*, 2019.
- Denil, M., Demiraj, A., and de Freitas, N. Extraction of salient sentences from labelled documents. *ArXiv*, abs/1412.6815, 2014.
- Devlin, J., Chang, M.-W., Lee, K., and Toutanova, K. Bert: Pre-training of deep bidirectional transformers for language understanding. *ArXiv*, abs/1810.04805, 2019.
- Dosovitskiy, A., Beyer, L., Kolesnikov, A., Weissenborn, D., Zhai, X., Unterthiner, T., Dehghani, M., Minderer, M., Heigold, G., Gelly, S., Uszkoreit, J., and Houshy, N. An image is worth 16x16 words: Transformers for image recognition at scale. *ArXiv*, abs/2010.11929, 2021.
- Fu, R., Hu, Q., Dong, X., Guo, Y., Gao, Y., and Li, B. Axiom-based grad-cam: Towards accurate visualization and explanation of cnns. *ArXiv*, abs/2008.02312, 2020.
- Gao, J., Wang, Z., Xuan, J., and Fidler, S. Beyond fixed grid: Learning geometric image representation with a deformable grid. *ArXiv*, abs/2008.09269, 2020.
- He, K., Chen, X., Xie, S., Li, Y., Doll’ar, P., and Girshick, R. B. Masked autoencoders are scalable vision learners. *ArXiv*, abs/2111.06377, 2021.
- Jiang, P.-T., Zhang, C.-B., Hou, Q., Cheng, M.-M., and Wei, Y. Layercam: Exploring hierarchical class activation maps for localization. *IEEE Transactions on Image Processing*, 30:5875–5888, 2021.
- Kanehira, A., Takemoto, K., Inayoshi, S., and Harada, T. Multimodal explanations by predicting counterfactuality in videos, 2019.
- Kaya, H., Gürpınar, F., and Salah, A. A. Multi-modal score fusion and decision trees for explainable automatic job candidate screening from video cvs. *2017 IEEE Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pp. 1651–1659, 2017.
- Kim, Y. Convolutional neural networks for sentence classification. In *EMNLP*, 2014.
- Lee, G., Jeong, J., Seo, S., Kim, C., and Kang, P. Sentiment classification with word attention based on weakly supervised learning with a convolutional neural network. *ArXiv*, abs/1709.09885, 2018.
- Li, J., Chen, X., Hovy, E. H., and Jurafsky, D. Visualizing and understanding neural models in nlp. In *HLT-NAACL*, 2016.
- Lu, J., Batra, D., Parikh, D., and Lee, S. Vilbert: Pretraining task-agnostic visiolinguistic representations for vision-and-language tasks. In Wallach, H. M.,

- Larochelle, H., Beygelzimer, A., d'Alché Buc, F., Fox, E. B., and Garnett, R. (eds.), *NeurIPS*, pp. 13–23, 2019. URL <http://dblp.uni-trier.de/db/conf/nips/nips2019.html#LuBPL19>.
- Miao, Y., Gowayyed, M. A., and Metze, F. Eesen: End-to-end speech recognition using deep rnn models and wfst-based decoding. *2015 IEEE Workshop on Automatic Speech Recognition and Understanding (ASRU)*, pp. 167–174, 2015.
- Muhammad, M. B. and Yeasin, M. Eigen-cam: Class activation map using principal components. *2020 International Joint Conference on Neural Networks (IJCNN)*, pp. 1–7, 2020.
- Nguyen, D. Comparing automatic and human evaluation of local explanations for text classification. In *NAACL*, 2018.
- Rao, K., Sak, H., and Prabhavalkar, R. Exploring architectures, data and units for streaming end-to-end speech recognition with rnn-transducer. In *2017 IEEE Automatic Speech Recognition and Understanding Workshop (ASRU)*, pp. 193–199, 2017. doi: 10.1109/ASRU.2017.8268935.
- Ribeiro, M. T., Singh, S., and Guestrin, C. "why should I trust you?": Explaining the predictions of any classifier. *CoRR*, abs/1602.04938, 2016. URL <http://arxiv.org/abs/1602.04938>.
- Rogers, A., Kovaleva, O., and Rumshisky, A. A primer in bertology: What we know about how bert works. *Transactions of the Association for Computational Linguistics*, 8:842–866, 2020.
- Selvaraju, R. R., Das, A., Vedantam, R., Cogswell, M., Parikh, D., and Batra, D. Grad-cam: Why did you say that? visual explanations from deep networks via gradient-based localization. *CoRR*, abs/1610.02391, 2016. URL <http://arxiv.org/abs/1610.02391>.
- Shrikumar, A., Greenside, P., and Kundaje, A. Learning important features through propagating activation differences, 2019.
- Simonyan, K., Vedaldi, A., and Zisserman, A. Deep inside convolutional networks: Visualising image classification models and saliency maps. *CoRR*, abs/1312.6034, 2014.
- Tjoa, E. and Guan, C. A survey on explainable artificial intelligence (xai): Toward medical xai. *IEEE Transactions on Neural Networks and Learning Systems*, 32:4793–4813, 2021.
- Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. Attention is all you need, 2017a.
- Vaswani, A., Shazeer, N. M., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. Attention is all you need. *ArXiv*, abs/1706.03762, 2017b.
- Wang, J., Tuyls, J., Wallace, E., and Singh, S. Gradient-based analysis of nlp models is manipulable. *ArXiv*, abs/2010.05419, 2020.
- Wu, J. and Mooney, R. J. Faithful multimodal explanation for visual question answering. *CoRR*, abs/1809.02805, 2018. URL <http://arxiv.org/abs/1809.02805>.
- Zeiler, M. D. and Fergus, R. Visualizing and understanding convolutional networks, 2013.

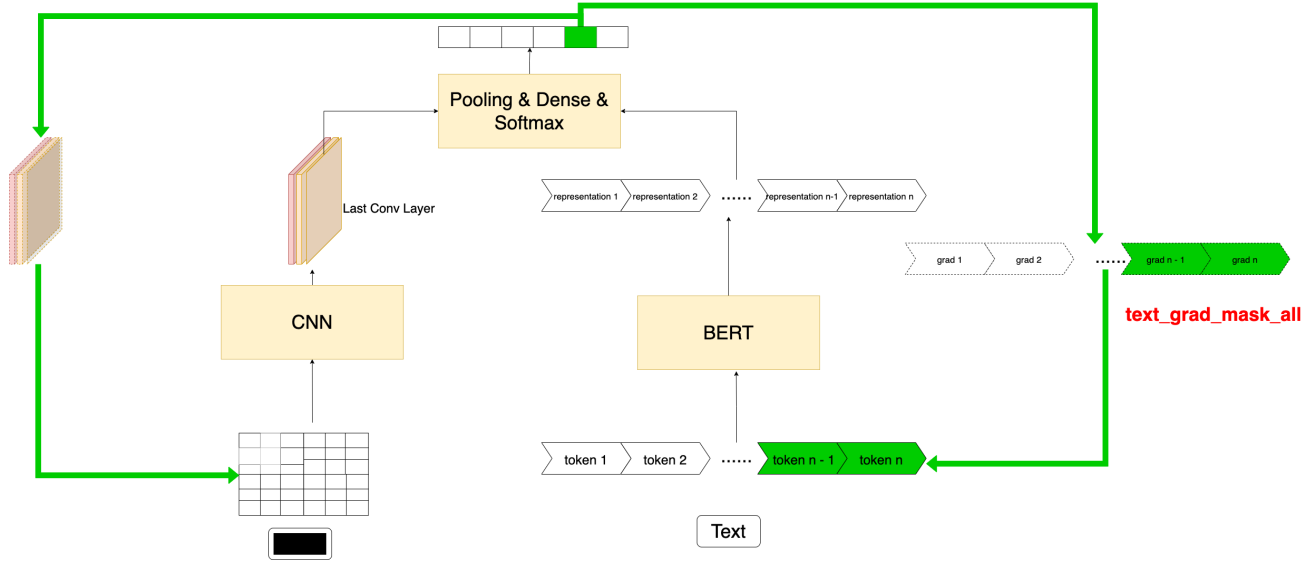


Figure 5. Mask the image completely, forward and backward pass only the text modality through the model and get text\_grad\_mask\_all.

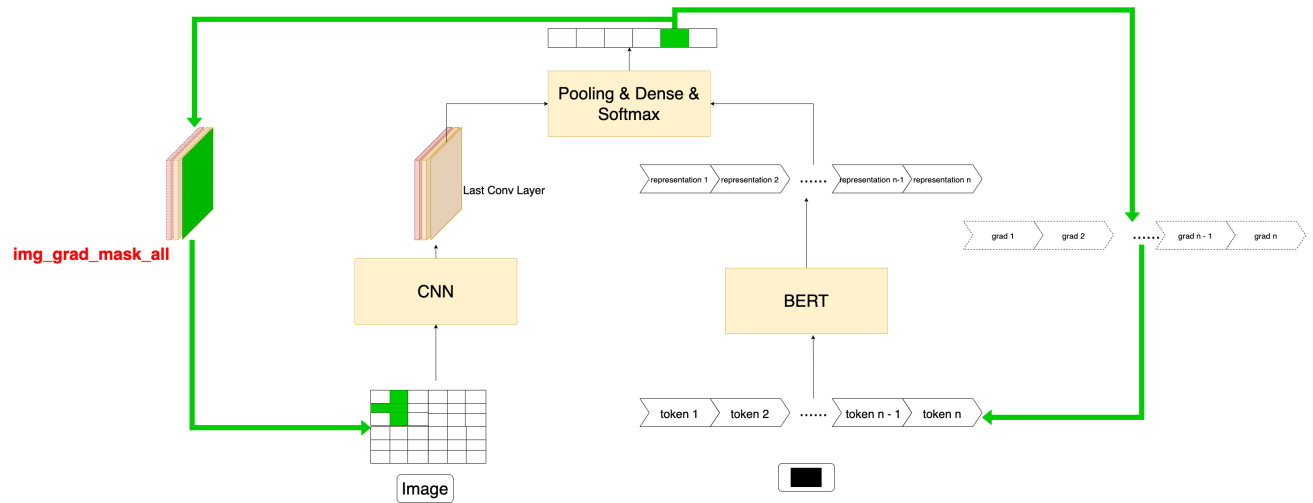


Figure 6. Mask the text completely, forward and backward pass only the image modality through the model and get img\_grad\_mask\_all.

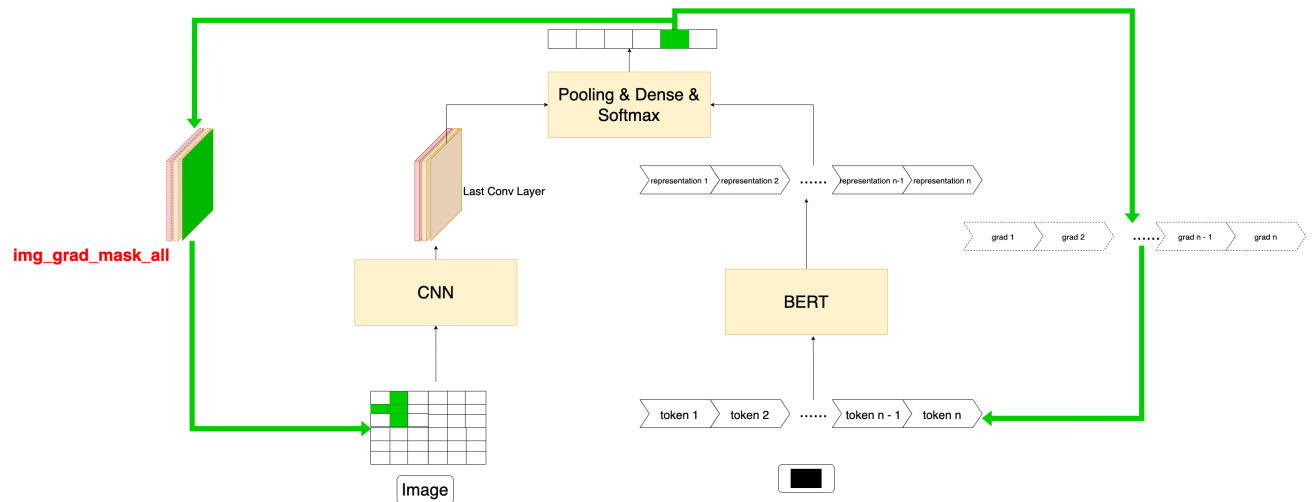


Figure 7. Mask the text completely, forward and backward pass only the image modality through the model and get img\_grad\_mask\_all.